

# Introduction to linear regression

(Week 05 lecture notes)

Wei Q. Deng (Presented by Tianle Chen)

Department of Statistical Sciences

August 1st 2018

# Topics covered in this lecture

- Machine learning the probabilistic way
- Supervised learning problems
- Regression
- Case studies in R and Python

# Why and what is machine learning ?

- We are in the era of **big data**: data of massive size and complex nature
  - 40 billion indexed web pages
  - the entire genome of thousands of individuals
  - more than 160 million credit card transactions per hour
- Machine learning provides **automated methods** to process and analyze data
  - automatically detect pattern in data
  - use the pattern to predict future data
  - use the pattern to inform decision making
- Many machine learning approaches rely on tools of probability theory

# Types of machine learning

- Supervised learning (or predictive learning)
  - classification (output is categorical)
  - regression (output is continuous)
- Unsupervised learning (or descriptive learning)
  - clustering
  - dimension reduction
  - graph structure
  - matrix completion
- Reinforcement learning (learning to make decisions)

# A general setting for supervised learning problems

Suppose we have some **inputs**  $X$  and some **outputs**  $y$ , the goal is to learn a mapping using a set of labelled set of input-output pairs  $\{(X_i, y_i)\}_{i=1}^n$ :

- also called a **training** dataset

To verify the quality of this mapping, we often use

- a **testing** data set containing (sometimes absent) just the **inputs**  $X_{\text{test}}$

The established mapping provides

- a model that describes the relationship between the inputs and outputs
- a model that can be used to predict outputs in the testing data

# Some examples of classification problems

When the outputs are categorical, these are known as the **classification** or **pattern recognition** problems:

- Handwritten digit recognition: postal offices process your letters automatically.
- How likely are you to get a new credit card approved?
- Identifying faces from images.

# handwritten digit recognition



# handwritten digit recognition

- the data consist of
  - images of handwritten digits from 0 to 9
  - labels identify by a human
- the goal is to process images automatically without humans hand labelling
- the learning algorithm need to achieve “near-human” performance



# credit card approval

- given attributes such as
  - age
  - past credit history
  - current financial status
  - number of credit cards owned
  - household income, etc.
- Credit card applications either accepted (+) or rejected (-)

## credit card approval

- Data can be found here
- Some analysis in R here
- Some analysis in Python here

# Some examples of regression problems

When the outputs are continuous or ordinal, these are known as the **regression** problems:

- How much can your neighbour's house go for in a market like this?
- Predict the trading value of the TSX index tomorrow or in a month given the current market condition and other available information.
- Predict the age of a current subscriber of the Globe and Mail based on their browsing history.
- Predict the amount of pollutant in a given postal code area with weather data, traffic condition and information on industrial activities in the area and nearby.

# house price prediction



**497 CLINTON ST**  
Toronto, Ontario M6G2Z3

**\$1,699,900**


5 + 2 3

Single Family House

MLS\* Number: C4199526



SUTTON GROUP-ADMIRAL REALTY INC.  
Brokerage



**#1704 -80 CUMBERLAND ST N**  
Toronto, Ontario M5R3V1

**\$1,099,000**

2 2

Single Family Apartment

MLS\* Number: C4198625



BAY STREET GROUP INC.  
Brokerage



**#PH 8 -102 BLOOR ST W**  
Toronto, Ontario M5S1M8

**\$2,250,000**

3 3

Single Family Apartment

MLS\* Number: C4198562



HOMELIFE/DIAMONDS REALTY INC.  
Brokerage



**#508 -181 DAVENPORT RD**  
Toronto, Ontario M5R1J1

**\$769,000**

1 + 1 1

Single Family Apartment



**210 ROBERT ST**  
Toronto, Ontario M5S2K7

**\$1,559,000**

3 3

Single Family House



**29 SUSSEX AVE**  
Toronto, Ontario M5S1J6

**\$2,999,000**

4 2

Single Family House

# house price prediction

- the listing price of all other houses in your neighbourhood is public (realtor.ca)
- a number of factors influence the listing price in your neighbourhood
  - house types (detached, semi-detached, etc.)
  - living space (square feet)
  - number of bedrooms
  - number of bathrooms
  - distance to a good public school
  - mortgage policies
  - availability of inventory
  - overall desirability of the neighbourhood
  - other amenities nearby
- How much should your neighbour list their house for in this market?

# Model TSX index

## Market Summary > S&P/TSX Composite Index INDEXTSI: OSPTX

[+ Follow](#)

**16,433.86** **-21.87 (0.13%)** ↓

Jul. 27, 12:03 p.m. EDT · Disclaimer

1 day

5 days

1 month

**1 year**

5 years

Max



Open  
High

16,465.36  
16,489.06

Low

16,430.61

# Model TSX index

- Index Characteristics such as constituents
- a number of economic factors influence the index
  - labor costs
  - interest rates
  - government policy
  - taxes
  - etc.
- recent news or economic developments
- Can you provide a one-day or one-month forecast for the TSX composite index?

How to approach these problems?




# What does data look like?

For the simplest case, we look at a dataset with one input vector.

i	X	Y
1	3.2	6.95
2	1.5	5.22
3	2.4	6.46
4	3.2	7.03
5	5.4	9.71
6	5.5	9.67
7	6.1	12.69
8	5.7	10.85
9	2.8	5.21
10	3.9	7.82

For example, the third observation is  $(X_3, Y_3) = (2.4, 6.46)$ . For real data, usually you don't have the index  $i$  column as given in the table.

# Visualize the data with a scatterplot



# Observing the scatterplot

- Most data seem to fall near a line
- For bigger values of  $X$ ,  $Y$  seems to be bigger as well
- In other words,  $X$  and  $Y$  seem to be positively related
- How can we find the line that “best” describes the relationship between  $X$  and  $Y$ ?

# A simple linear regression

- the output: an **outcome** variable  $Y$
- one input: a **predictor** variable  $X$
- We want to find a simple linear function of  $X = (x_1, \dots, x_n)$  that approximates  $Y = (y_1, \dots, y_n)$ :
  - A simple linear function has the form  $f(X) = a + bX$
  - “best” fit for a data point ( $i$ ) is measured by a small distance  $r_i = y_i - f(x_i)$
  - How do we aggregate the residuals ( $r_i$ ) to find the line of best fit?
- **Least squares method**

$$\operatorname{argmin}_{a,b} \sum_{i=1}^n r_i^2 = \operatorname{argmin}_{a,b} \sum_{i=1}^n (y_i - a - bx_i)^2$$

## Solving the least squares problem

$$\frac{\partial \sum_{i=1}^n r_i^2}{\partial a} = 2 \sum_{i=1}^n r_i \frac{\partial r_i}{\partial a} = 0$$

and

$$\frac{\partial \sum_{i=1}^n r_i^2}{\partial b} = 2 \sum_{i=1}^n r_i \frac{\partial r_i}{\partial b} = 0$$

Plug in  $r_i = y_i - f(x_i) = y_i - a - bx_i$  and solve the two equations simultaneously to obtain the least squares solutions.

$$\hat{a} = \bar{y} - \hat{b}\bar{x}$$

and

$$\hat{b} = \frac{\sum_{i=1}^n x_i y_i - \frac{1}{n} \sum_{i=1}^n x_i \sum_{i=1}^n y_i}{\sum_{i=1}^n x_i^2 - \frac{1}{n} (\sum_{i=1}^n x_i)^2}$$

# Regression coefficients

The output is approximated by a simple linear function has the form  
 $f(X) = a + bX$

- $a$  is the **intercept** (i.e. the value of the function when  $X = 0$ )
- $b$  is the **slope** (the strength of the positive or negative relationship between input and output)

The regression coefficients obtained using the least squares method:

- $\hat{a}$  is the estimated intercept
- $\hat{b}$  is the estimated slope

## Directions of association

- If the linear relationship is absent, then  $\hat{b}$  should be fairly close to 0
- If  $\hat{b} > 0$ , a **positive** relationship is possible
- If  $\hat{b} < 0$ , a **negative** relationship is possible

# Strength of association

- The strength of association between the input and output can be captured by **Pearson's correlation coefficient** *if* the input has only one variable (e.g. a simple linear regression model)
  - the correlation coefficient in a sample is denoted by  $r \in [-1, 1]$
  - the closer  $|r|$  is to 1, the stronger the relationship
  - between  $(-0.2, 0.2)$
- The strength of association between the input and output can be captured by **Pearson's correlation coefficient** *if* the input has only one variable (e.g. a simple linear regression model)
  - usually we square it and use  $r^2$
  - interpreted as the proportion of variance in  $Y$  that is explained by  $X$



# Regression in statistical software programs

The objective of the remaining lecture is to solve real prediction problems using statistical programming. You should be able to:

- Setup the regression problem
- Produce initial visualization of data
- Obtain an estimated linear regression model in either R or Python
- Read the outputs from R or Python
- Make prediction or interpret the results



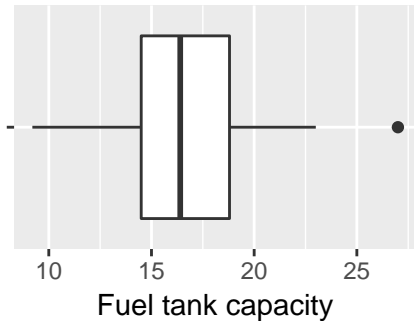
# Case study of regression in R and Python

- Data “Cars93” from the MASS library (appeared in lecture 1)
- Use data to answer the question:
  - Does car with a bigger fuel tank capacity necessarily have more horsepower?
  - For a car with fuel tank capacity of 27 US gallons, what is the maximum horsepower roughly?
  - What about a car with 9 US gallons?
  - What about a car with 16 gallons?




# About the dataset

- Data from 93 Cars on Sale in the USA in 1993
- Input variable is “Fuel tank capacity” in US gallon
- Output variable is “Horsepower” (maximum horsepower)
- Randomly choose 10 cars/observations to be the testing data
- The remaining 83 cars will be used to build the linear model





## Data visualization (R)




## Data visualization (R)



# Data visualization (Python)



# Data visualization (Python)



# Summarizing the data (R)

```
mean(Cars93[["Fuel.tank.capacity"]])
```

```
## [1] 16.66452
```

```
median(Cars93[["Fuel.tank.capacity"]])
```

```
## [1] 16.4
```

```
sd(Cars93[["Fuel.tank.capacity"]])
```

```
## [1] 3.27937
```

```
mean(Cars93[["Horsepower"]])
```

```
## [1] 143.828
```

```
median(Cars93[["Horsepower"]])
```

```
## [1] 140
```

```
sd(Cars93[["Horsepower"]])
```

```
## [1] 52.37441
```

```
round(cor(Cars93[["Fuel.tank.capacity"]], Cars93[["Horsepower"]]), 3)
```

```
## [1] 0.712
```



# Summarizing the data (Python)

```
import pandas
```

```
## /Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/importlib/_bootstrap.py:2  
##     return f(*args, **kwargs)
```

```
Cars93 = pandas.read_csv("Cars93.csv")  
import scipy  
from scipy.stats.stats import pearsonr
```

```
import pandas  
Cars93 = pandas.read_csv("Cars93.csv")  
Cars93[["Horsepower"]].mean()  
Cars93[["Horsepower"]].median()  
Cars93[["Horsepower"]].std()
```

```
## /Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/importlib/_bootstrap.py:2  
##     return f(*args, **kwargs)
```

```
Cars93[["Fuel.tank.capacity"]].mean()  
Cars93[["Fuel.tank.capacity"]].median()  
Cars93[["Fuel.tank.capacity"]].std()
```

```
import pandas  
Cars93 = pandas.read_csv("Cars93.csv")  
from scipy.stats.stats import pearsonr  
pearsonr(Cars93[["Fuel.tank.capacity"]], Cars93[["Horsepower"]])
```

## Estimated linear model in R

```
summary(lm(Horsepower~Fuel.tank.capacity, data = Cars93))

##
## Call:
## lm(formula = Horsepower ~ Fuel.tank.capacity, data = Cars93)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -96.321 -21.915  -5.349   14.863  120.528
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -45.613     19.968  -2.284   0.0247 *
## Fuel.tank.capacity   11.368       1.176   9.667 1.27e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 36.99 on 91 degrees of freedom
## Multiple R-squared:  0.5066, Adjusted R-squared:  0.5012
## F-statistic: 93.45 on 1 and 91 DF,  p-value: 1.268e-15
```

# Estimated linear model in Python

```
import pandas
import numpy as np
import stats
Cars93 = pandas.read_csv("Cars93.csv")
x = np.array(Cars93[["Fuel.tank.capacity"]])
y = np.array(Cars93[["Horsepower"]])
x.reshape((len(x), 1))
y.reshape((len(y), 1))
from sklearn.linear_model import LinearRegression
```

```
## /Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/imp
##     return f(*args, **kwds)
## /Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/imp
##     return f(*args, **kwds)
```

```
regression_model = LinearRegression()
regression_model.fit(x, y)
```

```
## /Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/site
##     linalg.lstsq(X, y)
```

```
intercept = regression_model.intercept_[0]
```

# Interpreting the results

- The model is written as “Horsepower~Fuel.tank.capacity”
- **Regression Coefficients:**

Predictors	Estimate	Std. Error	t value	Pr(>
(Intercept)	-45.613	19.968	-2.284	0.0247 *
Fuel.tank.capacity	11.368	1.176	9.667	1.27e-15 ***

- **Significance codes:**

Extremely Strong	Strong	Moderate	Weak	No evidence against $H_0$
0 ***	0.001 **	0.01 *	0.05 .	0.1-1

## Interpreting the results - Cont'd

- Residual standard error: 36.99 on 91 degrees of freedom
- R-squared
  - Multiple R-squared: **0.5066**
  - Adjusted R-squared: 0.5012
- F-statistic: 93.45 on 1 and 91 DF,
- **p-value:** 1.268e-15

What additional information do you observe?

## Back to the question

Does car with a bigger fuel tank capacity (FTC) necessarily have more horsepower?

- There is extremely strong evidence of a car with larger FTC having more horsepower.
- For one US gallon increase in fuel tank capacity, the maximum horsepower increases by 11 on average.

## Prediction problems:

- For a car with fuel tank capacity of 27 US gallons, what is the maximum horsepower roughly?
- What about a car with 9 US gallons?
- What about a car with 16 gallons?

## Prediction in R

```
new <- data.frame("Fuel.tank.capacity" = c(27, 9, 16))
predict(lm(Horsepower~Fuel.tank.capacity, data = Cars93),
        newdata = new)
```

```
##           1           2           3
## 261.32084  56.69841 136.27380
```



# Prediction in Python


```
x_new = [[27], [9], [16]]  
y_predict = regression_model.predict(x_new)  
print(y_predict)
```

```
## [[261.32083648]  
## [ 56.69840591]  
## [136.27379557]]
```

## Interpret the results

- For a car with fuel tank capacity of 27 US gallons, the *predicted* maximum horsepower is roughly 261.
- For a car with fuel tank capacity of 9 US gallons, the *predicted* maximum horsepower is roughly 57
- For a car with fuel tank capacity of 16 US gallons, the *predicted* maximum horsepower is roughly 136

Are these results trustworthy?



# Next class

- Linear regression as a statistical model
- Classic assumptions of linear regression
- More examples of linear regression
- Tests for equality of means as a special case of a simple linear regression
- Extending to multiple predictors in R